NEWS SUMMARIZER & CONNECTOR

# Obsidian & the Agent Stack

How we turn a spreadsheet of 8,926 news articles into a living, navigable knowledge graph — and how Claude Code orchestrates the agents that build it.

- Data & scripts
- Agent stack
- Obsidian vault

THE TOOL

# What is Obsidian?

A local-first knowledge base that stores everything as plain Markdown files on your own disk — no database, no lock-in, no cloud required.

### Plain Markdown

Every note is a `.md` file in a folder ("the vault"). Readable and editable with or without Obsidian.

### Wikilinks

Type `[[Canada]]` and notes link to each other. Links are bidirectional — every note shows its **backlinks**.

### Graph view

Obsidian renders all notes and links as an interactive graph. Densely-linked topics cluster together visually.

**Why it fits this project:** a news corpus is really a web of *themes*. Markdown + wikilinks + the graph turn 8,926 flat rows into something you can *see* and *walk*.

THE BIG PICTURE

# Spreadsheet → Agents → Knowledge graph

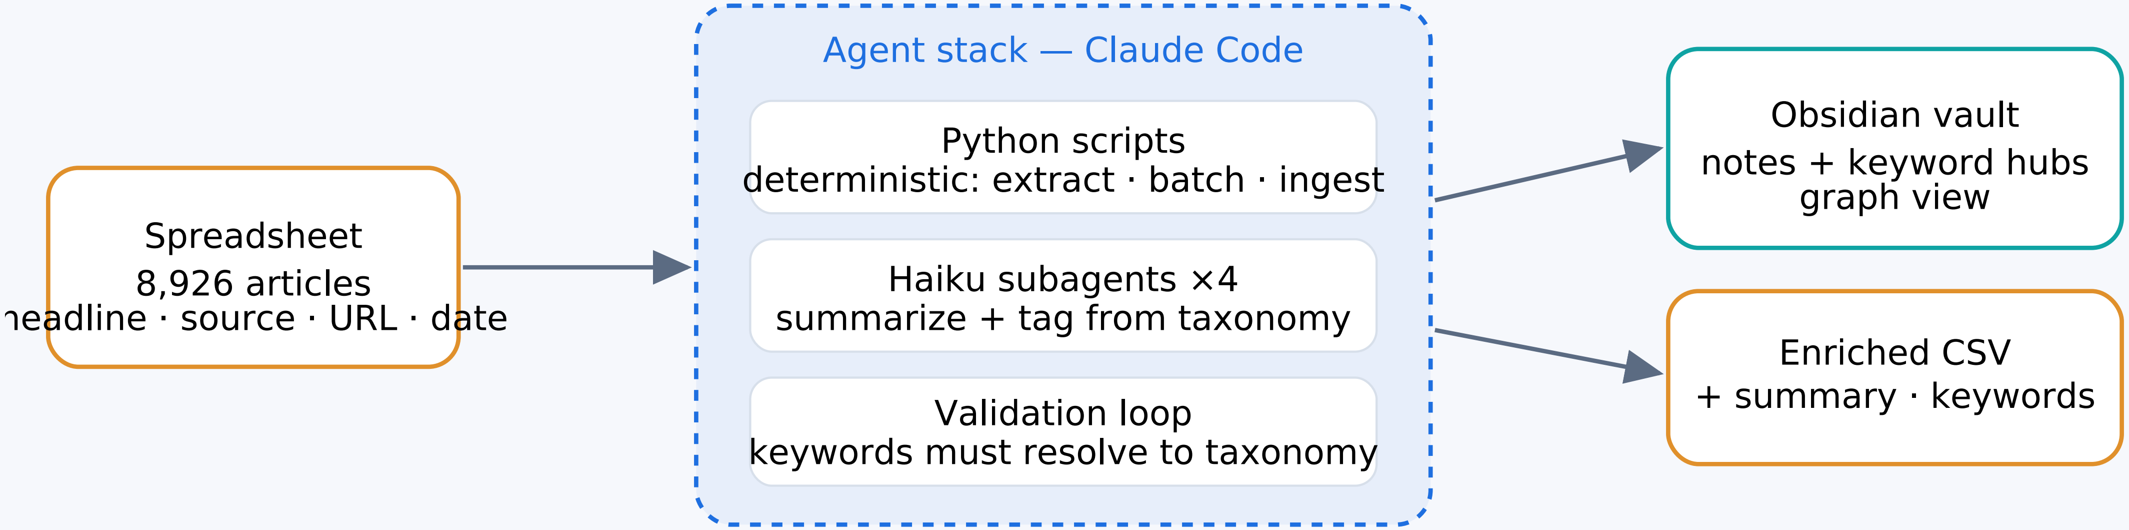


Deterministic plumbing in Python; judgement (summaries & tags) delegated to fast Haiku subagents; everything written as Markdown Obsidian reads natively. **Runs entirely on a Claude subscription — no API key.**

HOW WE USE IT

# Anatomy of the vault

Two kinds of note, linked by wikilinks:

- **Article notes** — one per article, in `vault/articles/news/`. YAML frontmatter (id, source, url, keywords, entities) + a 1–2 sentence **summary**, a one-line **key takeaway**, and a row of `[[Keyword]]` links.
- **Keyword hubs** — one per taxonomy term, in `vault/keywords/`. Every article that cites the keyword shows up as a **backlink**, so the hub becomes a theme dashboard for free.

62 keyword hubs · 100 article notes in the pilot · both grow as the pipeline runs.

```
--- frontmatter ---
id: n-0b11b0e5d85b
type: news-article
source: "Beta Kit"
keywords: ["Canada", "AI Adoption",
  "AI Policy & Regulation"]
entities: ["Cohere"]
---

## Summary
Cohere signs an exploratory AI
agreement with the Québec government.

## Key Takeaway
Signals a push to position Canada
as an AI hub.

## Keywords
[[Canada]] · [[AI Adoption]] ·
[[AI Policy & Regulation]]
```

HOW WE USE IT

# Wikilinks weave the graph

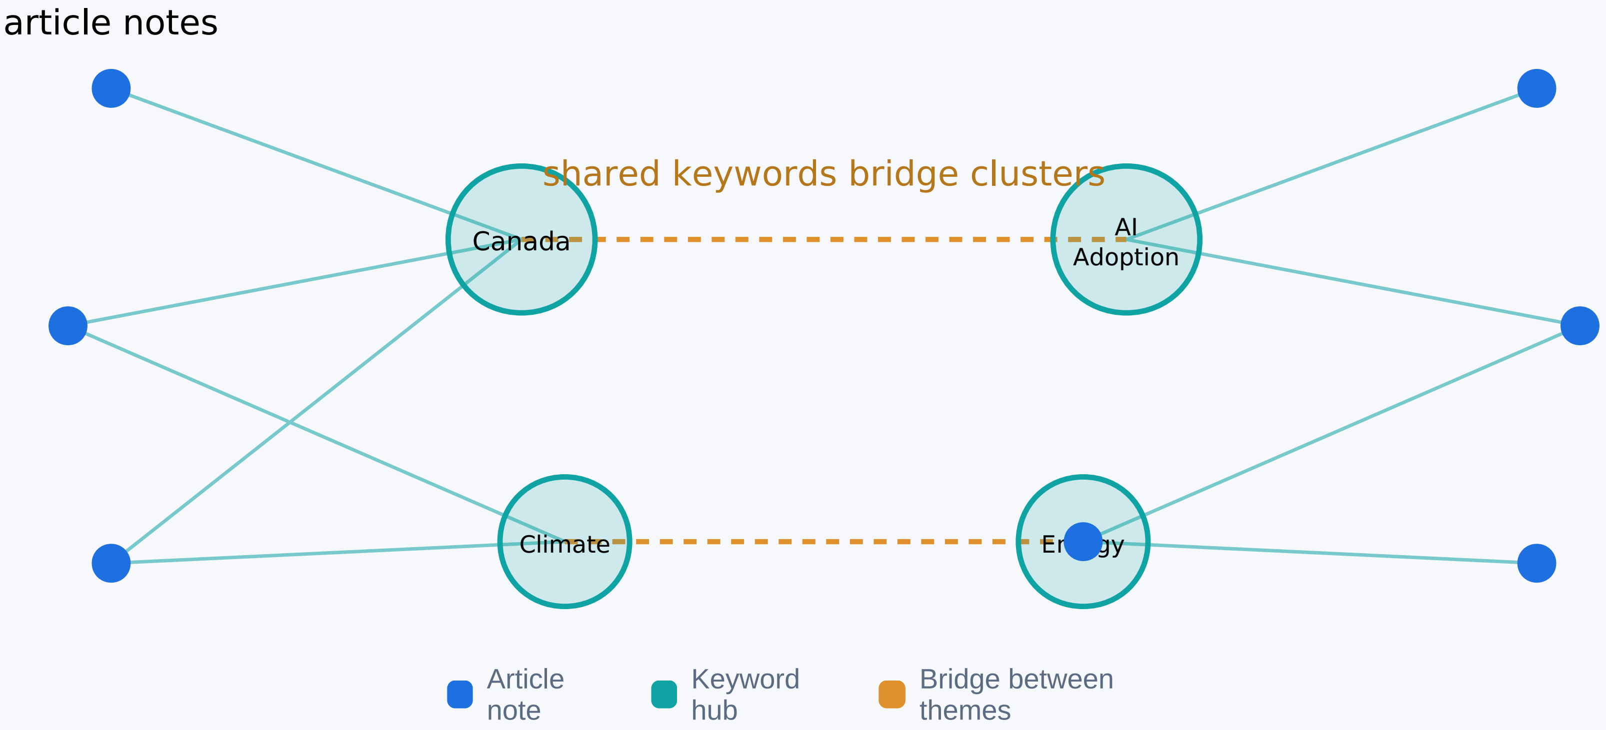

# Who does what

### Claude Code — orchestrator

Runs on your subscription. Drives the skills, spawns subagents, parses results, handles retries. The conductor, not the labourer.

### Skills — the playbooks

`/build-taxonomy` freezes the keyword vocabulary; `/process-articles N` runs the batch pipeline. Repeatable, resumable.

### Python scripts — determinism

Stdlib only. Extraction, stable IDs, batching, validation, note + CSV writing, graph rendering. No LLM guesswork where rules suffice.

### Haiku subagents — judgement

Up to 4 in parallel. Each reads a 25-article batch and returns summary, takeaway, keywords, entities. Cheap, fast, disposable.

**Division of labour:** rules-based work stays in Python (reproducible, free); only language understanding goes to the model — and to the *cheapest* model that does the job well.

WORKFLOW

# The processing pipeline, end to end

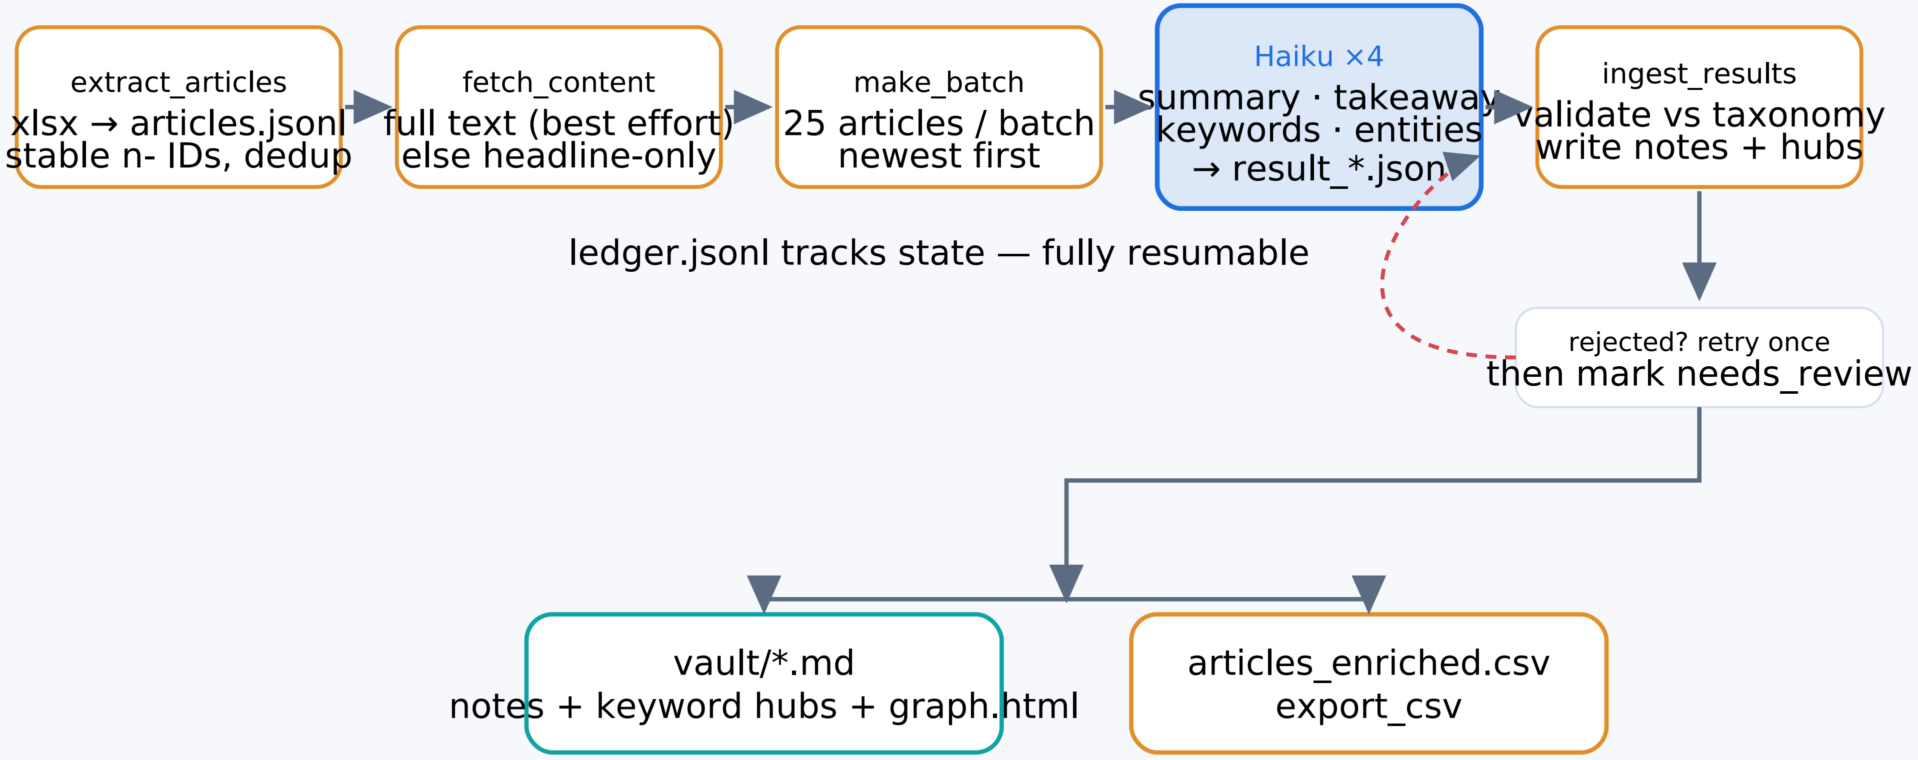

# Fan-out & the validation gate

The slow part — reading and summarizing — is parallelised:

- The orchestrator splits pending work into **25-article batches** and launches up to **4 Haiku subagents at once**.
- Each subagent must pick keywords **only from the frozen taxonomy**. Free-text tags are not allowed.
- `ingest_results.py` is the **gate**: keywords are resolved case-insensitively and alias-aware against the taxonomy. Anything that doesn't resolve is rejected.
- Rejected articles get **one retry** (with the reason fed back); a second failure is flagged `needs_review` — **the pipeline never wedges**.

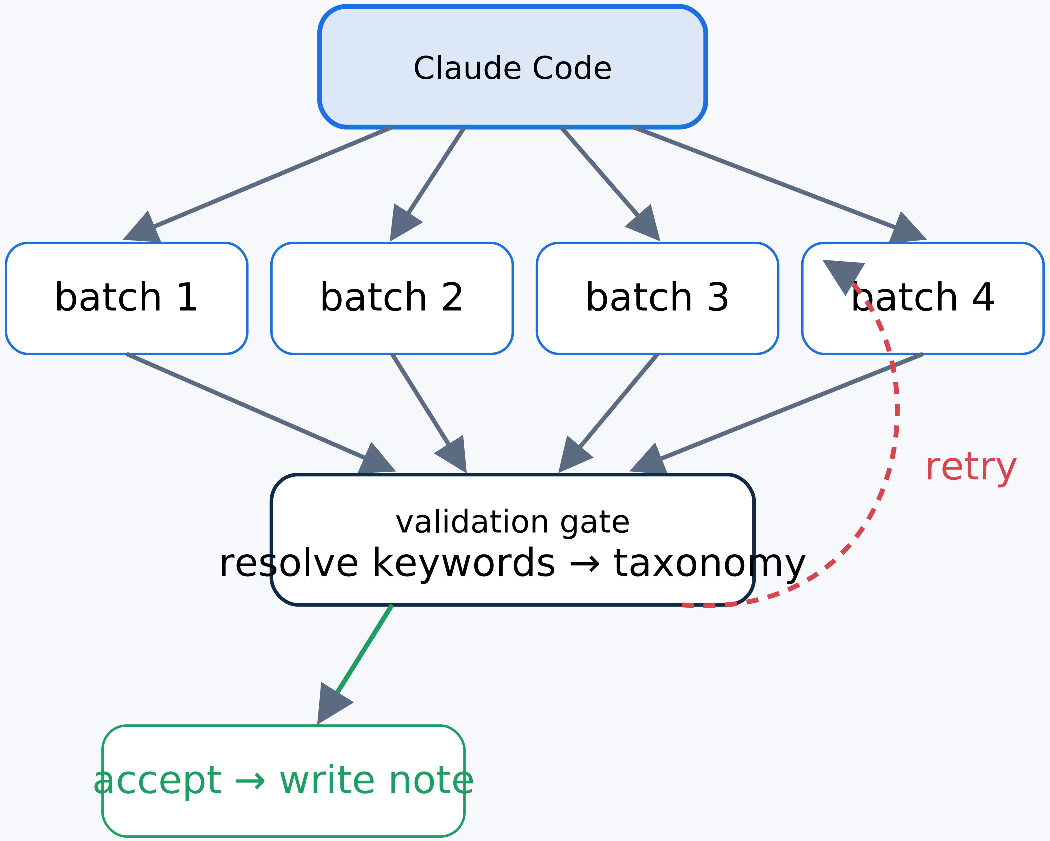

# The controlled taxonomy

~78 frozen keywords are the **shared vocabulary** that makes the graph cohere — and the bridge to the future insight-articles list.

- **Why controlled?** If every article invented its own tags, nothing would cluster. A fixed list of ~78 terms means thousands of articles snap onto the *same* hubs.
- **How built:** `/build-taxonomy` samples 400 stratified headlines, proposes grouped keywords for your review, then `finalize_taxonomy.py` freezes `taxonomy.json` (versioned).
- **Aliases** catch near-misses — `US / USA / America` all resolve to **United States**.

A slice of the vocabulary:

Canada · United States · Europe & UK · China · AI Adoption · AI Models & Research · AI Policy & Regulation · Startup Funding · Venture Capital · Big Tech · Climate Policy · Clean Energy · Data Centres · Cybersecurity · Economic Growth · + ~63 more

IN TANDEM

# How the agent stack & Obsidian work together

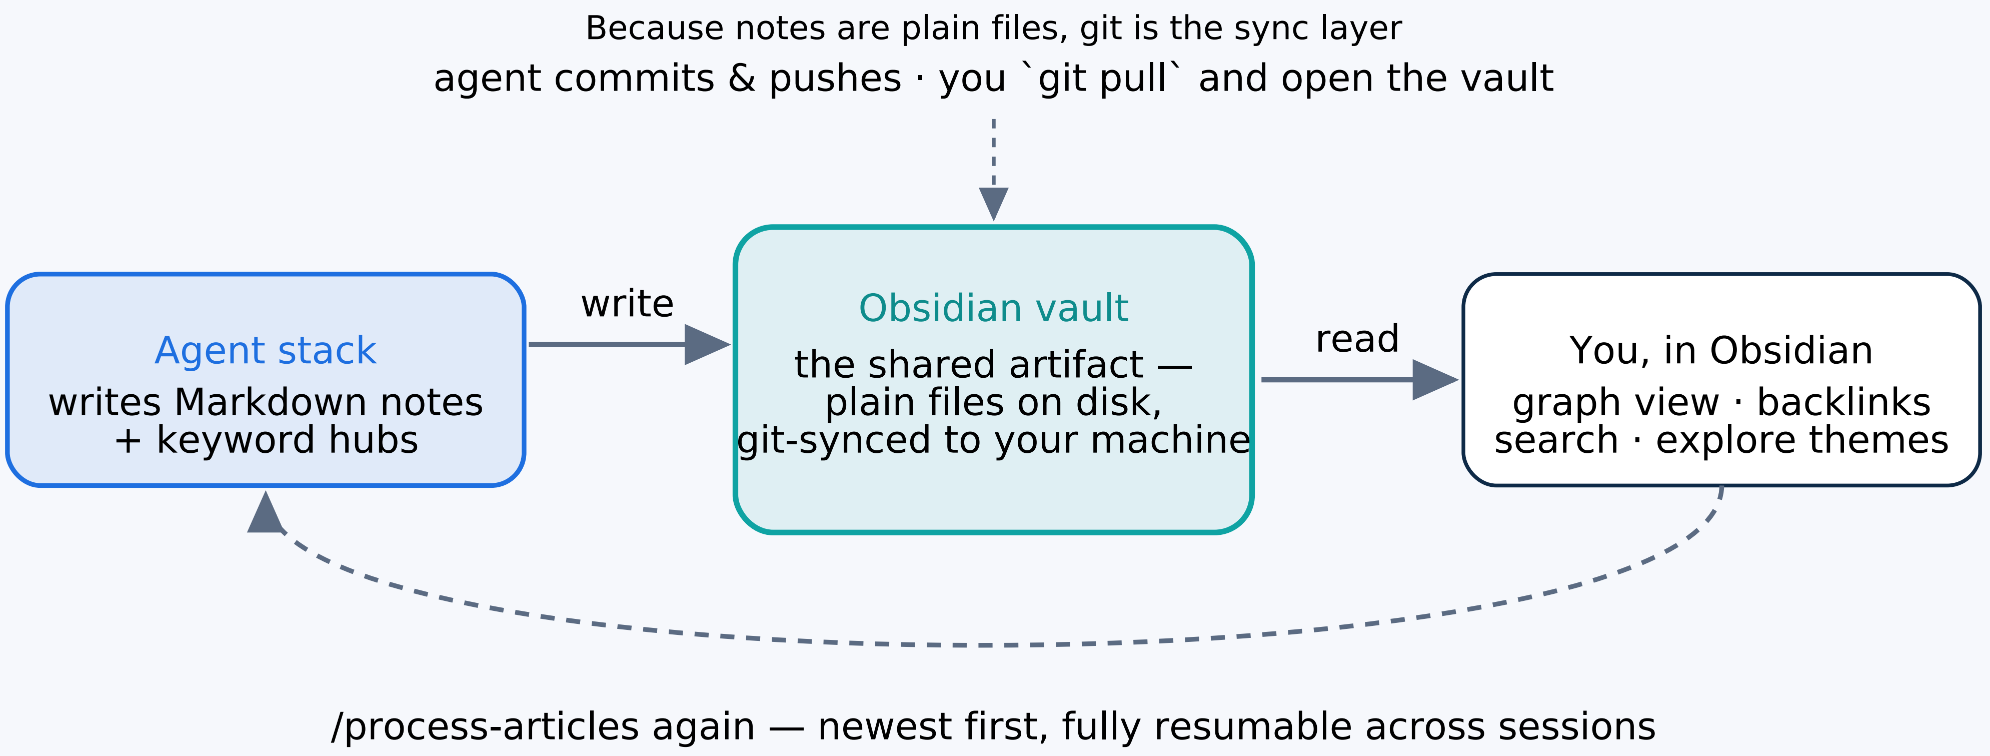

WHERE WE ARE

# Current state & next phase

**8,926**
articles in corpus

**100**
processed in pilot

**62**
keyword hubs live

### Built today

Extraction with stable dedup IDs · frozen 78-keyword taxonomy · Haiku batch summarization + tagging · Obsidian vault with graph · enriched CSV · standalone `graph.html`. Resumable via `ledger.jsonl`.

### Designed, next

Ingest the **insight-articles** list (`i-` IDs) tagged from the *same* taxonomy, then `find_connectors.py` ranks keywords by **(news × insight)** backlinks to surface the bridges between the two maps.

RECAP

# A spreadsheet becomes a map you can walk

### Scripts do the rules

Deterministic, free, reproducible.

### Haiku does the reading

Parallel, validated, cheap.

### Obsidian shows the shape

Local Markdown, graph, backlinks.

Open `vault/` as an Obsidian vault · view `output/graph.html` in any browser · run `/process-articles 100` to grow it.